

# Block-Diagonal LoRA for Eliminating Communication Overhead in Tensor aws Parallel LoRA Serving



Xinyu Wang



Jonas M. Kübler



Kailash Budhathoki



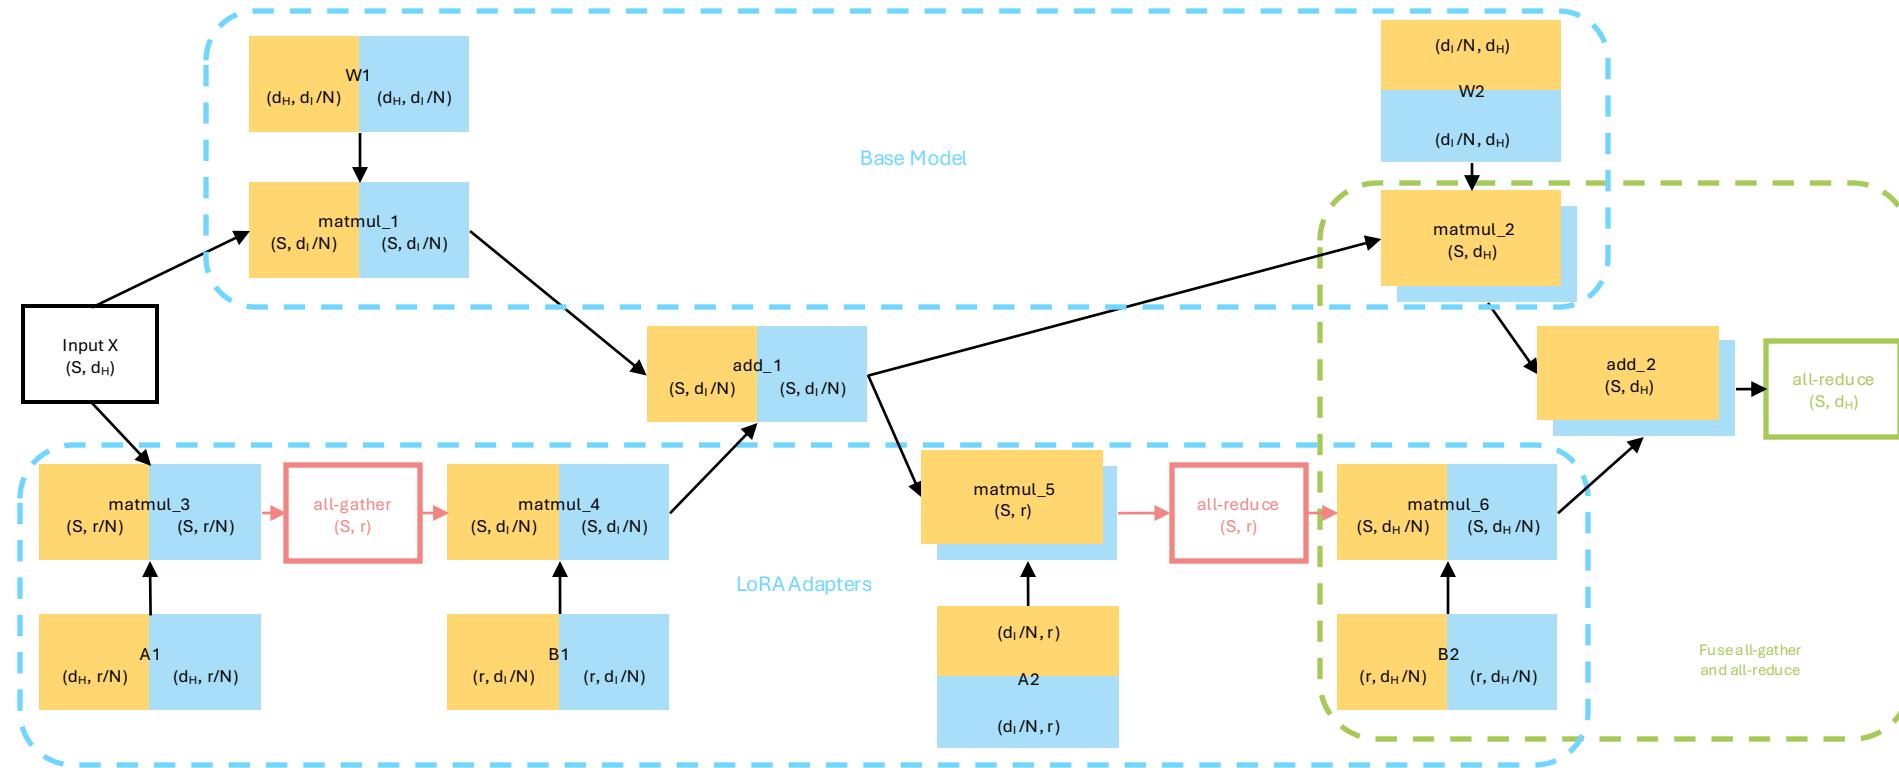
Yida Wang



Matthäus Kleindessner

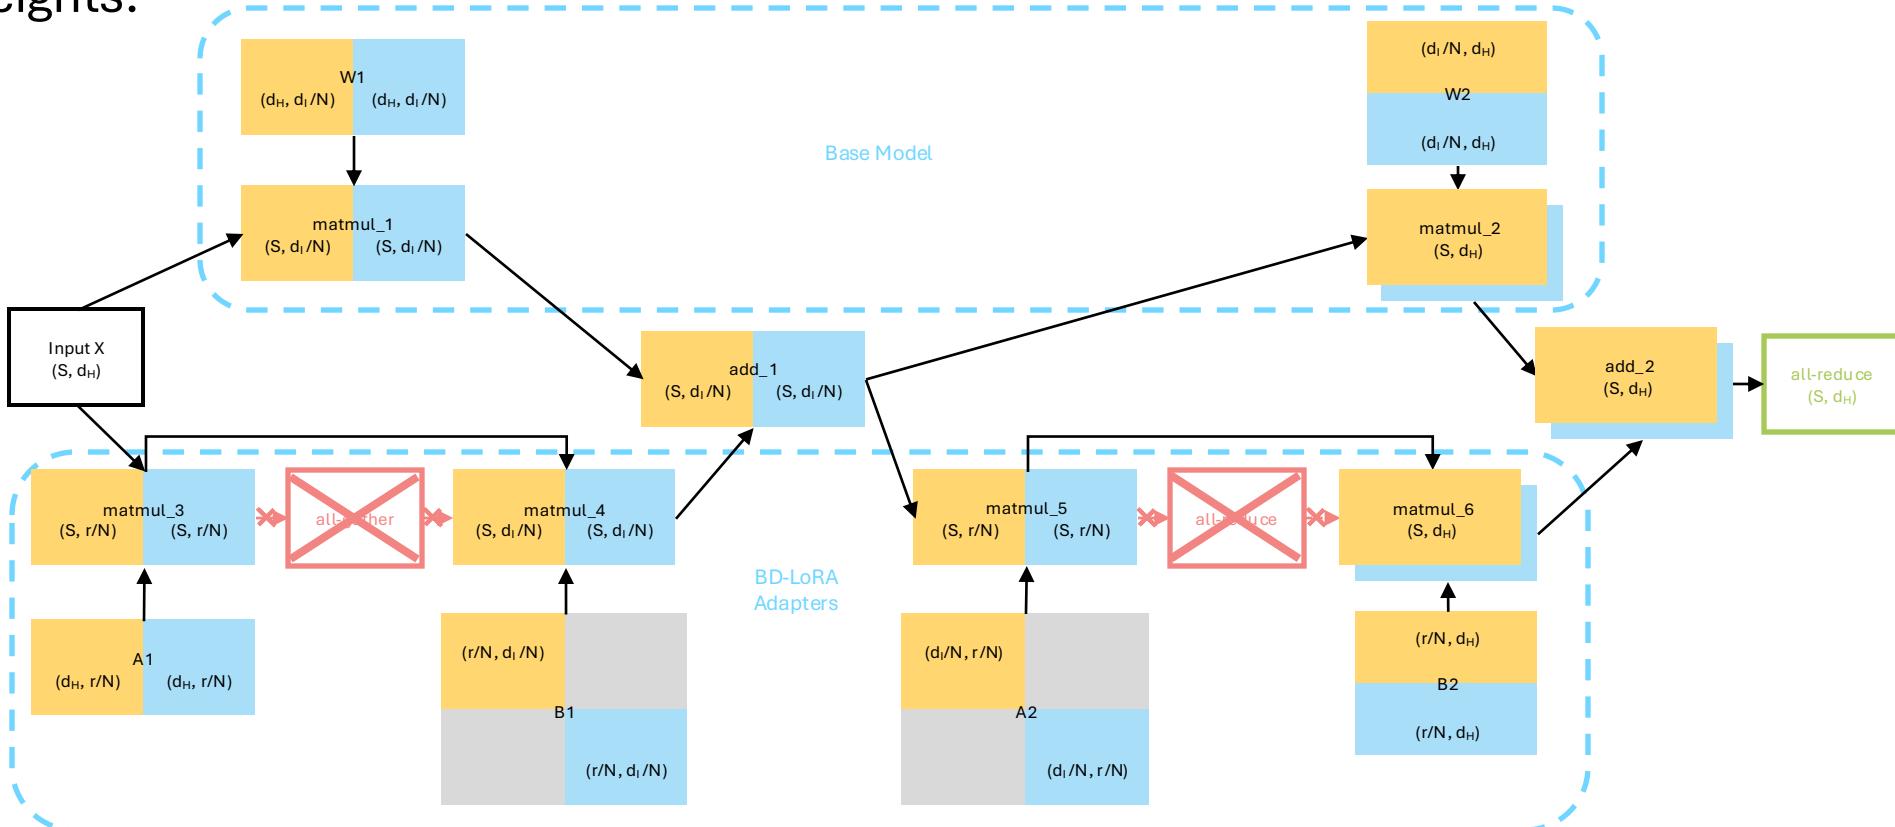
# Problem and Motivation

- For multi-LoRA serving with a shared base LLM, merging adapters prevents batching and incurs swap overhead.
- S-LoRA TP strategy adds extra all-gather/all-reduce beyond the base model.
- The additional cost is significant in practice and exacerbated at larger ranks.



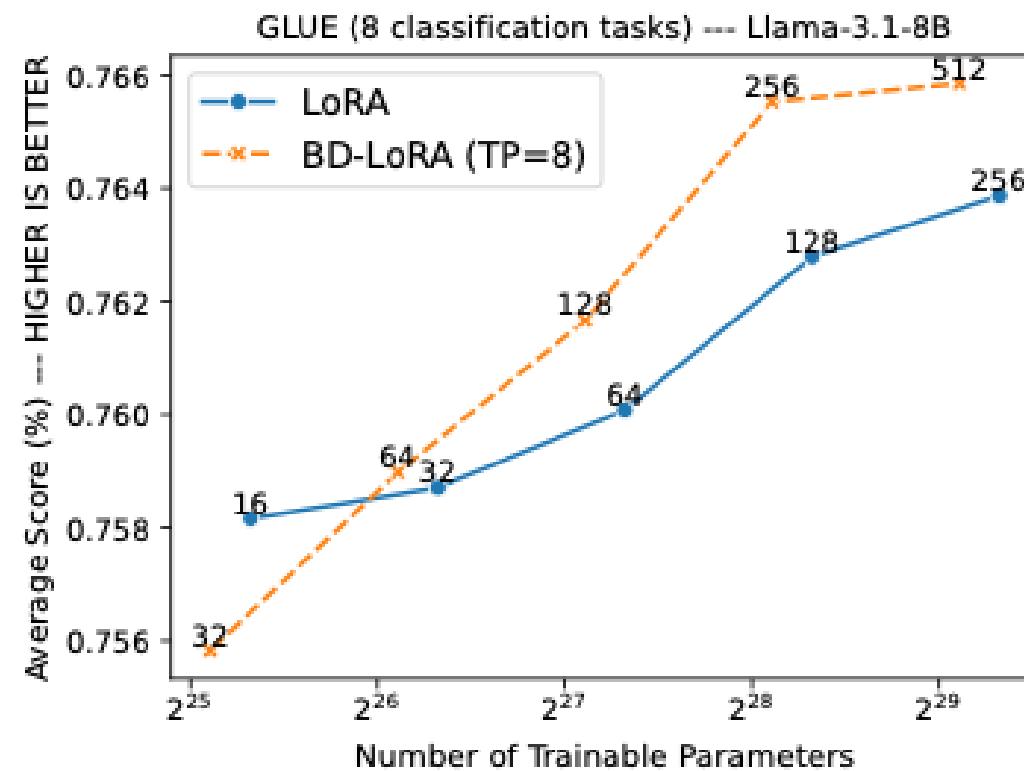
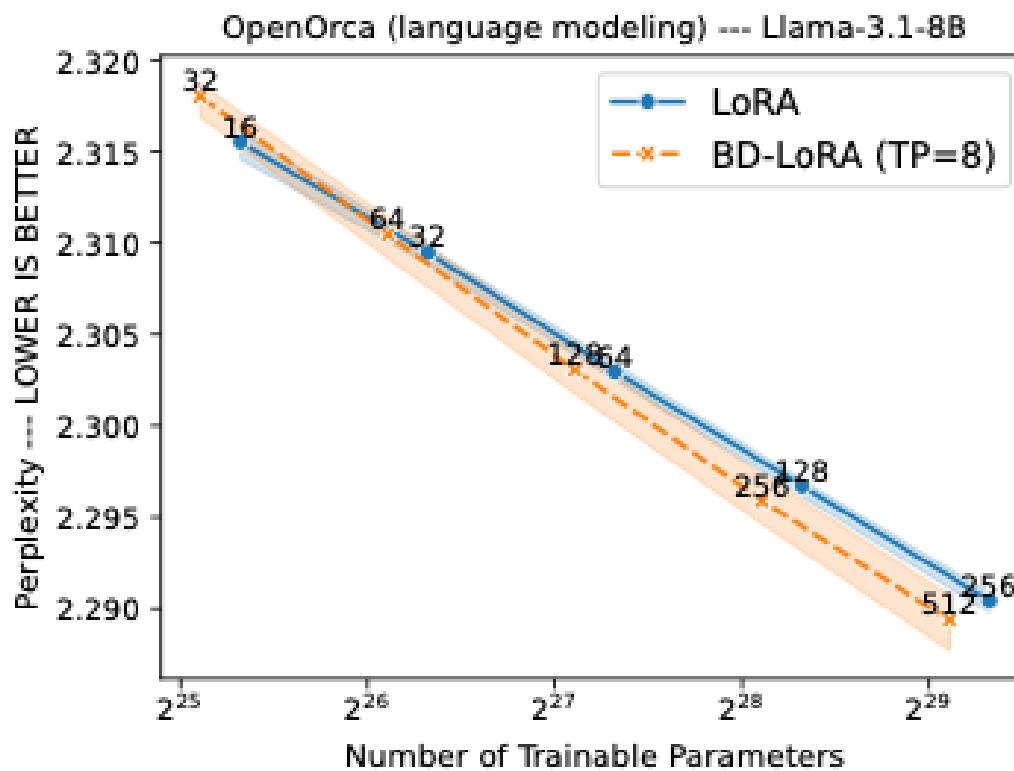
# BD-LoRA

- Constrain specific LoRA factors ( $B_1$  and  $A_2$ ) to be block-diagonal so they align with base TP sharding -> no extra all-gathers/all-reduces for LoRA.
- BD-LoRA can be interpreted as attaching independent LoRA adapters to each shard of the base weights.



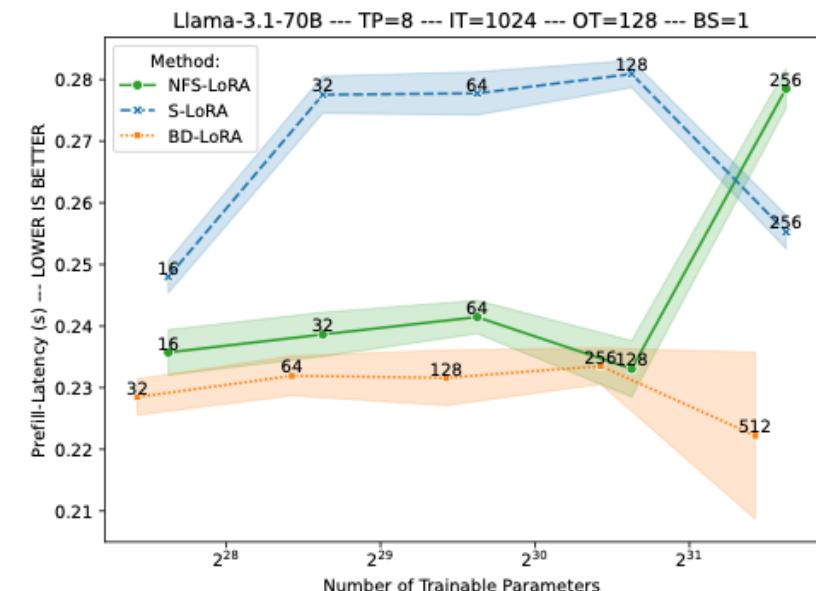
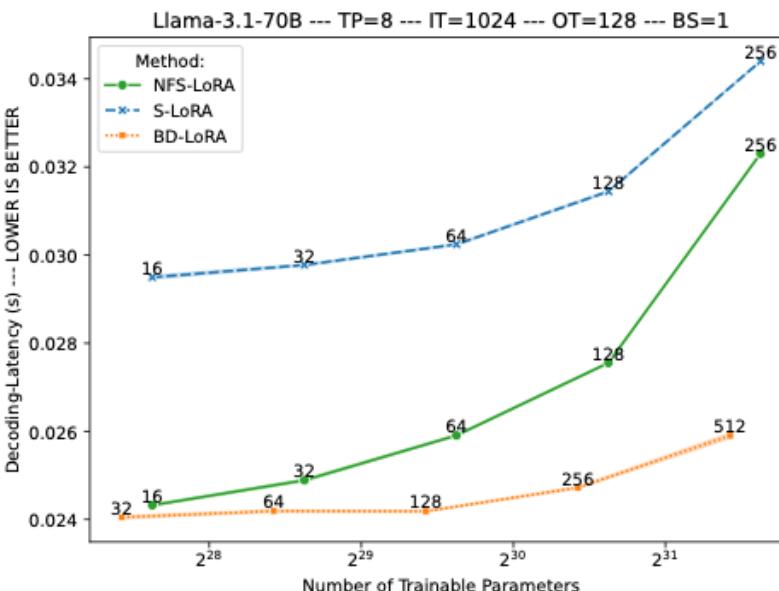
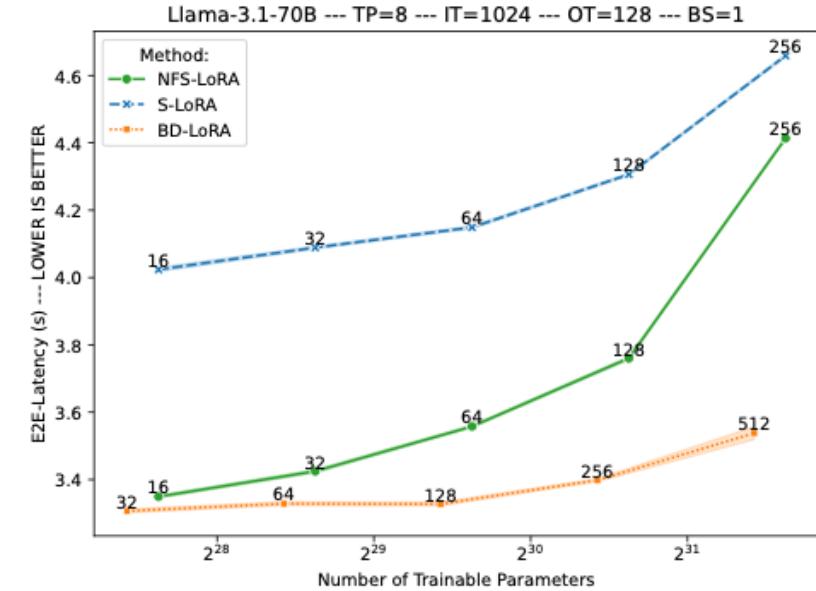
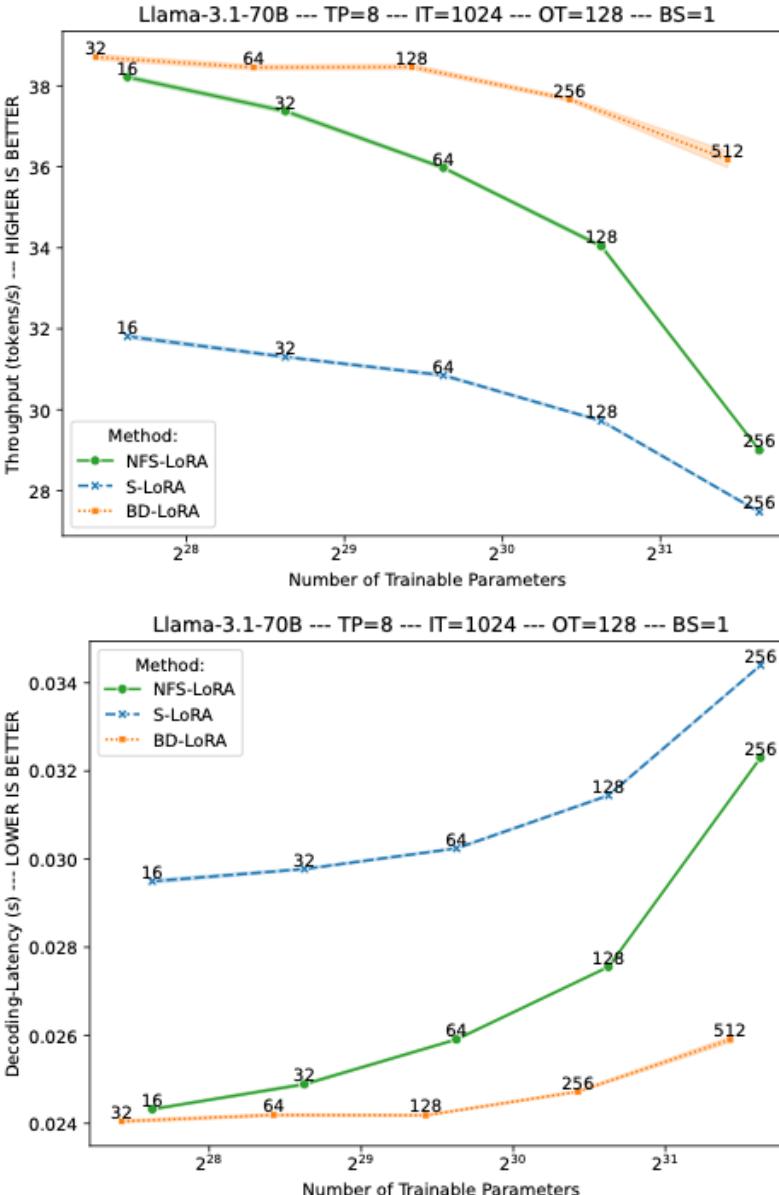
# Downstream performance

- For a similar number of trainable parameters, BD-LoRA shows comparable performance to LoRA.



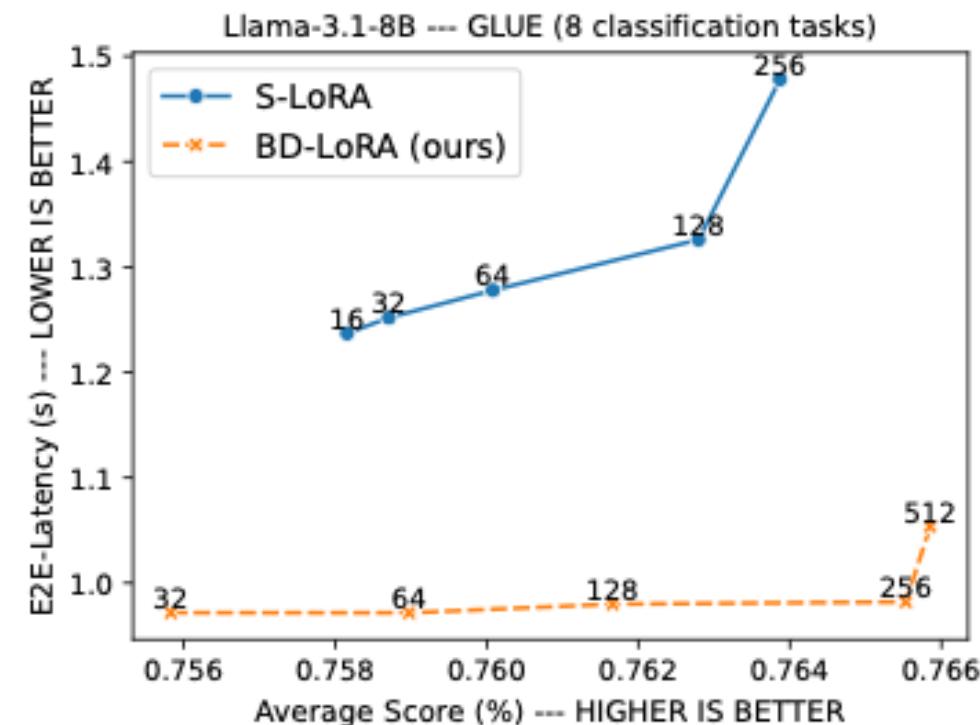
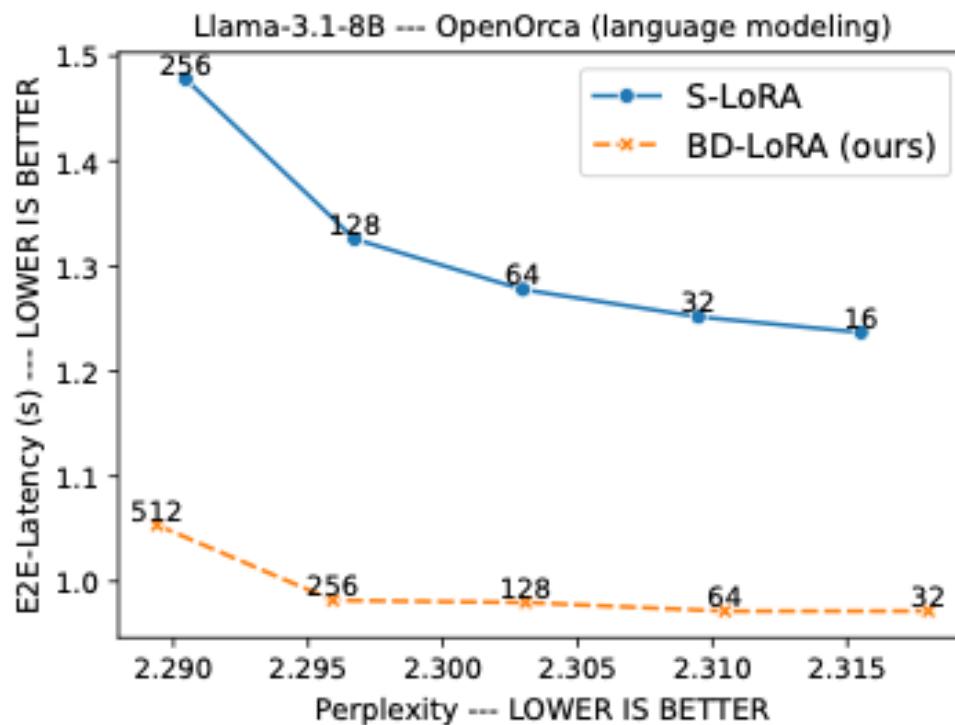
# Throughput and Latency

- BD-LoRA yields 1.79x (1.23x) end-to-end speed-up with 0.87x (1.74x) number of adapter parameters on 70B model and 1.63x (1.3x) with 0.86x (1.73x) parameters on 8B model.



# Downstream Performance vs. Run-time Trade-off

- At the same downstream performance, BD-LoRA achieves 1.27x – 1.51x end-to-end speedup.



# Takeaway

- BD-LoRA eliminates the additional communication overhead in S-LoRA.
- With the same number of trainable parameters, BD-LoRA and LoRA achieve similar downstream performance.
- With the same number of trainable parameters, BD-LoRA is faster than S-LoRA.
- With the same downstream performance, BD-LoRA is faster than S-LoRA.

Thank You!